Question 1

Not yet answered

Marked out of 1.00

Are micro-partitions user-configurable in Snowflake?

- a. No
- b. Yes

Question 2

Not yet answered

Marked out of 1.00

How does Snowflake determine which micro-partitions to scan during a query?

- a. Uses metadata filters based on pruning
- b. Applies machine learning
- c. Scans all micro-partitions
- d. Uses clustering keys

Question 3

Not yet answered

Marked out of 1.00

How does Snowflake handle changes in data distribution (e.g., skewed data)?

- a. Manual re-partitioning
- b. Auto-reclustering (with clustering keys)
- c. Requires data export and import
- d. Rewrites old partitions

Question 4

Not yet answered

Marked out of 1.00

Micro-partitions store data in which format?

- a. Row-based format
- b. Proprietary Snowflake log
- c. JSON
- d. Columnar format

Question 5

Not yet answered

Marked out of 1.00

What information does Snowflake store for each micro-partition?

- a. All of the above
- b. Data skew distribution
- c. Count of NULLs per column
- d. Min/Max values per column

Question 6

Not yet answered

Marked out of 1.00

What is a Micro-Partition in Snowflake?

- a. A block of storage used to store metadata only
- b. A user-defined partition of data
- c. A query optimization technique
- d. An automatically created contiguous storage unit

Question 7

Not yet answered

Marked out of 1.00

What is the advantage of smaller micro-partitions in Snowflake?

- a. More granular pruning and faster queries
- b. Better support for transactions
- c. Reduced storage cost
- d. Improved write performance

Question 8

Not yet answered

Marked out of 1.00

What is the typical size range of a Snowflake micro-partition?

- a. 10 GB and above
- b. 100 MB to 1 GB
- c. 1 KB to 5 MB
- d. 1 MB to 10 MB (compressed)

Question 9

Not yet answered

Marked out of 1.00

What kind of data structure is used to store metadata about micro-partitions?

- a. CSV indexes
- b. Column statistics and ranges
- c. B-Trees
- d. JSON

Question 10

Not yet answered

Marked out of 1.00

When you insert new data into a table, how are micro-partitions affected?

- a. Existing partitions are overwritten
- b. All data is re-partitioned
- c. New micro-partitions are automatically created
- d. Partitions stay unchanged

Question 11

Not yet answered

Marked out of 1.00

Which of the following best describes “partition pruning” in Snowflake?

- a. Caching frequent partitions
- b. Rewriting partitions
- c. Dropping unused partitions
- d. Skipping micro-partitions that don't match query filters

Question 12

Not yet answered

Marked out of 1.00

Which of the following best describes the immutability of micro-partitions?

- a. They are mutable but updated in batches
- b. They are recreated on each insert
- c. They are deleted after every query
- d. They are read-only after creation

Question 13

Not yet answered

Marked out of 1.00

Which of the following can improve the effectiveness of micro-partition pruning?

- a. Using well-designed clustering keys
- b. Querying without WHERE clauses
- c. Writing to the same table continuously
- d. Using semi-structured data

Question 14

Not yet answered

Marked out of 1.00

Which of the following tools can help monitor micro-partition behavior in Snowflake?

- a. SYSTEM$CLUSTERING_INFORMATION function
- b. Query Profiler
- c. Information Schema
- d. Storage Usage Dashboard

Question 15

Not yet answered

Marked out of 1.00

Which Snowflake feature heavily relies on micro-partition metadata for optimization?

- a. Query Result Caching
- b. Materialized Views
- c. Automatic Clustering
- d. Failover regions